

Node.js Deployment Patterns

Node Live London, July 2016

Luke Bond



T: @lukeb0nd | E: luke@yld.io | G: @lukebond

What You Will Learn From This Talk

- Common Node.js architectures & how best to deploy them
- Process monitors (spoiler: use Linux)
- Scaling and service discovery considerations
- A little about containers along the way

Node.js Deployment Patterns


For simplicity, let's consider the following common app types:

- Simple web-app
- Simple web-app + API + websockets + cache + database
- Microservice-based distributed system with a message bus

Let's explore deployment patterns for these types of apps.

TL;DR: don't overcomplicate; you probably don't need Kubernetes.

Simple Web-App



Simple Web-App

- Easily horizontally scalable
- No communication between components or across hosts
 - Hence no service discovery needed
- Run one Node process for each core
- Scale boxes horizontally; every box is the same
- Load balance (e.g. *nginx*, *HAProxy*, *balance*)

Recommendation: use **systemd** (or **PM2** if you must!).

Aside: Process Monitors

There are various popular ones in use today for Node.js:

- mon
- nodemon
- forever
- PM2

The first three are simple process restarters; PM2 does has a built in load-balancer, and does logging, deployment and monitoring.

Use systemd as your Process Monitor


- Core PM2 features can be easily recreated with *systemd* & *balance*
- See my talk here where I did this in 20 minutes:
 - Video: <https://opbeat.com/events/nodeconf-oslo-2016/#deploying-and-running-node-js-to-production-in-2016>
 - Slides: <https://github.com/lukebond/nodeconf-oslo-20160604/blob/master/nodeconf-oslo-20160604.pdf>
- Read/watch this for the *how-to*; I won't go into it here

Why Bother when PM2 Just Works?

- The tooling is better (especially `journalctl`)
- You have more control; you can choose your components
- Run your Node.js apps the same way you run non-Node.js apps
- Any Linux sysadmin will understand it
- No need to reinvent the wheel
- Learn more about the OS on which your apps run

`</Aside>`

More Complex App w/ API + WS + DB



More Complex App w/ API + WS + DB


- Easily horizontally scalable
- Communication between components required
- Scale out WS & API by process and across hosts
 - Communication between hosts then required
 - Host affinity becomes important; e.g. Redis on same host as APIs
- How to do service discovery without too much complexity?

More Complex App w/ API + WS + DB

Recommendations: you need service discovery but keep it simple

- Use **AWS CloudFormation** and/or **Elastic Beanstalk**
- Use **Google App Engine**
- Use **Fleet** and a dynamically configured **nginx** or **HAProxy**
 - Easy to transition from using **systemd** to **CoreOS Fleet** 🙌
- Use the new **Docker Swarm** in Docker 1.12

Distributed Microservices System



Distributed Microservices System

- You don't want to be thinking about any plumbing
- You don't want to be specifying exactly what runs where
- You want a declarative platform with service discovery
- You want a platform that will reschedule on app or host failures

Distributed Microservices System

Recommendation:

- Use **Kubernetes**
- Use hosted Kubernetes such as **Tectonic** or **Google App Engine**
- Use a PaaS such as **Deis**
- **Docker Swarm** someday- not yet proven at scale IMO

Conclusion & General Guidelines

- **Keep it simple** or suffer the operational pain
 - No unnecessary complexity in stack or deployment
- Being **12-factor** saves you operational pain & brings portability
- **Containers** help if your application is sufficiently complex
 - Devs can test and debug whole system on their laptop
 - Homogeneous deployments even for heterogenous stacks
- Test, test, test; smoke test releases, ensure you can rollback easily
- **systemd** over **PM2** 💪

Thanks!



T: @lukeb0nd | E: luke@yld.io | G: @lukebond